

Benchmark paper

Anonymized Authors

Anonymized Affiliations
`email@anonymized.com`

Abstract. Diffusion MRI is widely used to explore brain microstructure, and numerous machine learning methods have been proposed to extract predictive information. Yet progress remains difficult to assess: studies rely on single datasets, heterogeneous pipelines, limited evaluation, and often lack accessible code, preventing meaningful comparison and reproducibility.


We introduce a modular and transparent benchmarking framework for machine learning in dMRI microstructure. The benchmark unifies multiple representative datasets within a standardized end-to-end pipeline and enables systematic comparison across tissue types, feature representations and models.

It reveals differences between white and gray matter analyses, quantify the impact of diffusion microstructure feature choices, and show that fast linear baselines can rival more complex deep learning methods. Standardized benchmarking is essential to move the field beyond proxy tasks and toward robust methods with real clinical relevance.

Keywords: Benchmark · Brain · diffusion MRI.

1 Introduction

1.1 Scientific and Clinical Motivation



- Diffusion MRI provides quantitative microstructural biomarkers
- Increasing use of ML for prediction (age, diagnosis, cognition, etc.)
- Lack of standardization in preprocessing and feature representation
- High variability in reported performance across studies

Table idea

Study	Preprocessing	Standardized?	Feature Type	Public Code?
Study A	Pipeline A	No	Voxel-wise	No
Study B	Pipeline B	No	ROI-based	Yes

1.2 Gap in the Field

- No standardized benchmark for diffusion MRI prediction
- Heterogeneous preprocessing pipelines
- Inconsistent feature representations (voxel-wise vs ROI-based), not analysis cross microstructure features
- Evaluation protocols not directly comparable
- Limited reproducibility and transparency, missing code
- Lack of results of the same experiments on other datasets

1.3 Contributions

- Standardized preprocessing pipeline for diffusion MRI
- Unified feature extraction strategy
- Controlled evaluation protocol (train/test splits)
- Comparison of multiple feature representations
- Public release of code and benchmark configuration
- End-to-end analysis
- Integration of both classical and DL models

2 Benchmark

* PLOT: Scheme of the benchmark

Our benchmark aggregates datasets with each dataset preprocessed with the same pipeline, with the exception of HCP which has its own preprocessing pipeline. We perform the feature extraction/data preparation and apply rigorously the analysis and cross validation that we are going to explain.

2.1 Datasets

- Number of subjects
- Age range / demographics
- Scanner details
- Acquisition protocol or summary table with acquisition parameters
- Prediction targets (classification/regression)

* PLOTS: Demographic distribution (age histograms) | Table with demographics summary (gender classes, age range...)

2.2 Preprocessing and Feature Extraction

Preprocessing (Himanshu)

- Denoising
- Motion and eddy correction
- Spatial normalization
- Masking strategy

Feature extraction

- White and Gray matter computation (Maybe merge in Matter representation and include a table with differences between both representations(?))
- Input: Microstructure information

Gray Matter Representation

- Voxel-wise features
- High-dimensional representation (64k features)
- Mask-based extraction

Explain that voxel-wise features preserve spatial resolution but increase dimensionality and risk of overfitting.

White Matter Representation

- Tract-based aggregation
- Mean values per tract (48 features)
- Reduced dimensionality
- Anatomical interpretability

Justify tract-averaged representation: Noise reduction, improved stability, lower dimensional feature space, anatomically meaningful structure

2.3 Prediction

- Feature normalization
 - Binary targets (encoded 0/1), Continuous targets (standardized during training)
 - Standardization fitted on training set only
 - Applied to validation/test sets
 - Inverse transform for regression outputs (for visualization)
- Task definitions and labels and reasons for task selection
 - Classification: Binary diagnosis prediction (e.g., ASD vs. TD)
 - Regression: Age prediction (brain age estimation)
 - Task selection rationale: Clinical relevance, data availability, and established benchmarks in the literature
- Evaluation Metrics
 - Classification: Accuracy, AUC, F1-score
 - Regression: R^2 , MAE
 - Mean $\pm$ std across folds
 - Statistical comparison tests (e.g., paired t-tests, Wilcoxon signed-rank tests) to assess significance of performance differences between models
- Models Used
 - Dummy baseline (e.g., majority class for classification, mean predictor for regression)
 - Classical ML pipeline: Linear models (e.g., logistic regression, linear regression), Random Forests, Support Vector Machines (SVM) + their PCA versions
 - Deep Learning models (e.g., dinov2, medicalnet)
 - Unified hyperparameter search strategy (e.g., grid search or random search with cross-validation on the training set)

2.4 Validation, Evaluation and Visualizations

- Data Splitting Strategy (Fixed train/test splits)
- Cross-validation within training set. Stratification
- No information leakage
- Explanation of plots and statistics

3 Results

3.1 Within-Dataset Performance

- Model/Feature/Tissue comparison
- Mean $\pm$ std of best run / model
- Statistical testing
- PLOT: Bar plots per dataset | Camcan-HCP comparison(?)

3.2 Feature Representation Comparison

- Gray vs white matter
- Microstructure metrics comparison
- PLOT: Tissue performance according to microstructure metric for task purposes | Microstructure features

3.3 "Handmade" vs DL embeddings

3.4 Training Size Analysis

* PLOT: Learning Curve

3.5 Key Findings

- Which models perform best? Which feature representations perform best? Tissue predominance in task prediction?
- Are simpler models sufficient? Are DL models necessary for predicting demographics?
- Stability across folds?

3.6 Methodological Insights

- Importance of dimensionality control
- Effect of tract-based aggregation
- Risk of overfitting in mesh representations without PCA bcs of the high number of vertices representations

4 Conclusion

- Introduced standardized dMRI benchmark for age prediction in healthy individuals using three large datasets (HCP, CamCAN, ABIDE)
- Compared feature representations (MD, SH, MK, RTOP) and models (dummy classifier, linear, `pca_linear`, forest, `pca_forest`, svm, `pca_svm`)
- Provided reproducible end-to-end framework and enables fair model comparison. Reduces methodological heterogeneity
- A step toward methodological unification/code availability

4.1 Limitations/Further steps

- No multi-site harmonization (explicitly as scope limitation). Would training in one dataset and predicting in another dataset have the same effect as predicting in different folds of the same dataset? Are the predictions very far?
- Limited tasks / ranges in prediction (e.g. HCP age prediction)
- Limited preprocessing / including other preprocessing parameters or corrections

For citations of references, we prefer the use of square brackets and consecutive numbers. Citations using labels or the author/year convention are also acceptable. The following bibliography provides a sample reference list with entries for journal articles [1], an LNCS chapter [2], a book [3], proceedings without editors [4], and a homepage [5]. Multiple citations are grouped [1–3], [1, 3–5].

Acknowledgments. A bold run-in heading in small font size at the end of the paper is used for general acknowledgments, for example: This study was funded by X (grant number Y).

Disclosure of Interests. It is now necessary to declare any competing interests or to specifically state that the authors have no competing interests. Please place the statement with a bold run-in heading in small font size beneath the (optional) acknowledgments¹, for example: The authors have no competing interests to declare that are relevant to the content of this article. Or: Author A has received research grants from Company W. Author B has received a speaker honorarium from Company X and owns stock in Company Y. Author C is a member of committee Z.

References

1. Author, F.: Article title. *Journal* **2**(5), 99–110 (2016)
2. Author, F., Author, S.: Title of a proceedings paper. In: Editor, F., Editor, S. (eds.) *CONFERENCE 2016, LNCS*, vol. 9999, pp. 1–13. Springer, Heidelberg (2016). <https://doi.org/10.1007/1234567890>

¹ If EquinOCS, our proceedings submission system, is used, then the disclaimer can be provided directly in the system.

3. Author, F., Author, S., Author, T.: Book title. 2nd edn. Publisher, Location (1999)
4. Author, A.-B.: Contribution title. In: 9th International Proceedings on Proceedings, pp. 1–2. Publisher, Location (2010)
5. LNCS Homepage, <http://www.springer.com/lncs>, last accessed 2023/10/25